

# LLM fine-tuning for time series annotation

Philippe Helluy

IRMA, Université de Strasbourg, Inria Tonus

August 29, 2025

# Outlines

How it works (roughly...)

Fine tuning of Qwen LLM for time series annotation

How it works (roughly...)

## Very brief history

- ▶ Artificial neural networks have existed for a **long time** ([Rosenblatt, 1958], late 1950s).
- ▶ Ups and downs, then **Yann LeCun** ([LeCun *et al.*, 1989], handwriting recognition, 1989).
- ▶ *Attention is all you need* [Vaswani *et al.*, 2017], invention of **transformers** at Google.
- ▶ Without huge **computing power**, it would not work.

## Completion is all you need

- ▶ Principle: we are given the beginning of a text. The task is to predict the next word. Example: "the cat eats the ..." (we must guess "mouse").

- ▶ Words (or *tokens*):

$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6 = t_m$
.	cat	the	eats	tomcat	mouse

- ▶ Corpus: "the cat eats the mouse.", "the tomcat eats the mouse.", "..the cat eats.", "..the mouse eats.", "..the tomcat eats.", etc.
- ▶ Remark: all sentences have  $\ell = 6$  words (we pad with the filler word ".").

# Digitization

- ▶ Encoding: to each word (or *token*) we associate a vector of  $m = 6$  dimensions

$$\text{"."} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \text{"cat"} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \text{"the"} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \text{etc.}$$

- ▶ Embedding into a lower-dimensional space of size  $p$  (to account for synonyms, among other things). For example  $p = 5$ . The embedding  $E_{w_0}$  is a function from  $\mathbb{R}^m$  to  $\mathbb{R}^p$ .
- ▶ Each token  $t_i$  is represented by a vector  $v_i$ . The embedding parameters  $w_0$  are unknown.

$$v_i = E_{w_0}(t_i).$$

# Encoder

- ▶ A sentence is thus represented by a "tensor" of  $\ell$  numerical vectors stacked together:

$$r_0 = v_{i_1} v_{i_2} \dots v_{i_\ell}$$

It is therefore an object in a space of dimension  $N = \ell \times p = 30$ .

- ▶ The sentence goes through  $k$  layers of transformers  $T_{w_i}$ , which are mappings from  $\mathbb{R}^N$  to  $\mathbb{R}^N$  with unknown parameter vectors  $w_i$

$$r_i = T_{w_i}(r_{i-1}), \quad i = 1 \dots k.$$


- ▶ The deeper we go into the layers, the more abstract the representation  $r_i$  of the initial sentence becomes. The vector  $r_k$  ("latent state") contains the information extracted by the network from the initial sentence  $r_0$ .

# Decoder

Finally, the decoder allows us to predict a probability vector  $p$  in  $\mathbb{R}^m$ :  $p_i$  is the probability that the next word is  $m_i$ .

$$p = D_{w_{k+1}}(r_k).$$

In summary (more details in [Tunstall *et al.*, 2022, Vigon, 2023]):





# Training

- ▶ Choosing the form of the functions  $E_{w_0}$ ,  $T_{w_i}$ ,  $D_{w_{k+1}}$ : a trade-off between cost, efficiency, and simplicity. For now, it is as much an art as it is a science.
- ▶ Historically, several possible architectures: CNN, RNN, LSTM, transformers. Evolution is strongly linked to available computing power.
- ▶ The parameter vector  $w$ , of size  $s$ , is **unknown**.
- ▶ Training consists in optimizing these parameters so that the model best reproduces the sentences from the corpus.
- ▶ This is the most difficult part of the computation: it requires a supercomputer, specialized processors, and costs millions of euros.
- ▶ Orders of magnitude for GPT-3:  $\ell = 2000$ ,  $m = 50000$ ,  $p = 20000$ ,  $s = 170$  billion...

# Inference

- ▶ Once the parameters  $w$  are computed, inference is fast.
- ▶ A network can be retrained for a specific task at a reduced cost (*fine tuning*). For example, Copilot and ChatGPT are specialized versions of GPT- $x$  ( $x = 3, 4, 5$ )
- ▶ The *pre-prompt* is essential to obtain high-quality results.
- ▶ For computational cost reasons, **ChatGPT does not learn in real time**. Its short-term memory is therefore limited to a few thousand words.

## Example with PyTorch

- ▶ PyTorch is a free software library developed by Huggingface to manipulate neural networks.
- ▶ Minimal example for creating a toy LLM from scratch with PyTorch (ask me).

# Fine tuning of Qwen LLM for time series annotation

# Objectives

- ▶ Teach a small LLM (Qwen2.5 0.5 billions parameters) to describe a time series
- ▶ Example of prompt:

Describe the time series in three sentences.

First sentence: describe trend (increasing/decreasing/flat).

Second sentence: noise intensity (low/medium/high).

Third sentence: approximate localisation of global maximum (beginning/middle/end) and global minimum (beginning/middle/end).

Put the description in a JSON format with the following pattern

```
<json>{ "trend": <sentence1>,  
  "noise": <sentence2>,  
  "extrema": <sentence3> }  
</json>
```

Series: [02, 01, 00, 03, ... ]

- ▶ Tokenization is essential: the values in the series are scaled to the range 00..99.

# Necessity of training

- ▶ Small LLM fails to answer properly. It generates total garbage.

- ▶ Example of output:

```
series name: \"series\"
time interval: 1 day
``` To provide a detailed analysis of the given time series data, I will break it down into its components.

### Time Series Description:
The provided time series consists of daily values for several variables over a period of one year (2023-01-01 to 2023-12-31).

### Trend Analysis:
- **Increase**: The values increase steadily from 01 to 04, then decrease again.
- **Decrease**: The values decrease from 04 onwards
...
```

- ▶ Let's try to improve this with supervised fine tuning

## Practical methodology

- ▶ Generate a dataset of correct examples with a large LLM (Mistral, ChatGPT, etc.)
- ▶ Use 90% of the dataset for training and the 10% left for evaluation.
- ▶ Apply a supervised fine tuning (SFT) procedure on a small LLM from this dataset.
- ▶ In order to reduce the cost we adopt a LoRA approach. <sup>1</sup>
- ▶ Now go to [https://github.com/phelluy/DLAA\\_2025](https://github.com/phelluy/DLAA_2025) and follow the README file.
- ▶ The Colab GPU memory is limited, I had to reduce the size of the computations. For full results see [Boileau *et al.*, 2025].

---

<sup>1</sup>The LoRA (Low-Rank Adaptation) approach in supervised fine-tuning (SFT) freezes the original model weights and injects small trainable low-rank matrices into certain layers (typically linear projections in attention/FFN). This drastically reduces the number of parameters that need updating, making fine-tuning large models much more memory- and compute-efficient while still achieving strong adaptation.

# Bibliographie I



BOILEAU, M., HELLUY, P., PAWLUS, J. et VYETRENKO, S. (2025).

Towards interpretable time series foundation models.

*arXiv preprint arXiv:2507.07439.*



LECUN, Y., BOSER, B., DENKER, J. S., HENDERSON, D., HOWARD, R. E., HUBBARD, W. et JACKEL, L. D. (1989).

Backpropagation applied to handwritten zip code recognition.

*Neural computation*, 1(4):541–551.



ROSENBLATT, F. (1958).

The perceptron: a probabilistic model for information storage and organization in the brain.

*Psychological review*, 65(6):386.



TUNSTALL, L., VON WERRA, L. et WOLF, T. (2022).

*Natural language processing with transformers.*

" O'Reilly Media, Inc. ".



VASWANI, A., SHAZEER, N., PARMAR, N., USZKOREIT, J., JONES, L., GOMEZ, A. N., KAISER, Ł. et POLOSUKHIN, I. (2017).

Attention is all you need.

*Advances in neural information processing systems*, 30.



VIGON, V. (2023).

Cours sur les réseaux de neurones.

<http://octaviogame.com/liens/>.